

# Functional Programming II: Writing Functions



UNIVERSITY OF  
SAN FRANCISCO

James D. Wilson

BSDS 100 - Intro to Data Science with R



All  $\mathbb{R}$  functions have three parts

- the `body()`, the code inside the function
- the `formals()`, the list of arguments which controls how you can call the function
- the `environment()`, the *map* of the location of the function's variables



```
> myFunc <- function(x) x^2
```

```
> myFunc  
function(x) x^2
```

```
> formals(myFunc)  
$x
```

```
> body(myFunc)  
x^2
```

```
> environment(myFunc)  
<environment: R_GlobalEnv>
```



- 1 Write a function that takes two arguments, `a` and `b`, and returns rows `a` through `b` of `mtcars`
- 2 Write a function that takes a numeric vector as an input, squares every value in the vector, appends the squared vector to the original vector in the form of a data frame, and prints the first 10 rows of the data frame to the console
- 3 Write a function that takes a numeric vector as an input, squares every value in the vector, appends the squared vector to the original vector in the form of a data frame, and then returns and stores the data frame **over** the original vector, i.e., replace the old vector (which was input) with the new data frame (which is returned)



- Scoping is the set of rules that govern how  $\mathbb{R}$  looks up the value of a symbol
- There are four basic principles behind  $\mathbb{R}$ 's implementation of lexical scoping:
  - 1 name masking
  - 2 functions vs. variables
  - 3 a fresh start
  - 4 dynamic lookup



```
rm(list=ls())

myFunc_01 <- function() {
  x <- 1
  y <- 2
  c(x,y)
}

myFunc_01()
```

What does the preceding code return?

```
[1] 1 2
```

The function searches inside itself for `x` and `y`



```
rm(list=ls())

myFunc_01 <- function() {
  x <- 1
  y <- 2
  c(x,y)
}

myFunc_01()
```

What does the preceding code return?

```
[1] 1 2
```

The function searches inside itself for `x` and `y`



If a name isn't defined inside a function, R will look one level up

```
x <- 2

myFunc_02 <- function() {
  y <- 1
  c(x, y)
}

myFunc_02()
```

What does the preceding code return?

```
[1] 2 1
```

- What if you omitted `x <- 2`?





If a name isn't defined inside a function, R will look one level up

```
x <- 2
```

```
myFunc_02 <- function() {  
  y <- 1  
  c(x, y)  
}
```

```
myFunc_02()
```

What does the preceding code return?

```
[1] 2 1
```

- What if you omitted `x <- 2`?



```
rm(list=ls())

x <- 1
myFunc_03 <- function() {
  y <- 2
  myFunc_04 <- function() {
    z <- 3
    c(x,y,z)
  }
  myFunc_04()
}
```

What does the preceding code return?

```
[1] 1 2 3
```

Search begins inside the function, then where that function was



```
rm(list=ls())

x <- 1
myFunc_03 <- function() {
  y <- 2
  myFunc_04 <- function() {
    z <- 3
    c(x,y,z)
  }
  myFunc_04()
}
```

What does the preceding code return?

```
[1] 1 2 3
```

Search begins inside the function, then where that function was



```
rm(list=ls())

x <- 1
myFunc_03 <- function() {
  x <- 1000
  y <- 2
  myFunc_04 <- function() {
    x <- 99
    z <- 3
    c(x, y, z)
  }
  myFunc_04()
}
myFunc_03()
```

What does the preceding code return?

```
[1] 99 2 3
```



```
rm(list=ls())

x <- 1
myFunc_03 <- function() {
  x <- 1000
  y <- 2
  myFunc_04 <- function() {
    x <- 99
    z <- 3
    c(x, y, z)
  }
  myFunc_04()
}
myFunc_03()
```

What does the preceding code return?

```
[1] 99 2 3
```

# Functions vs. Variables



The same principles apply for finding functions just as they do for finding variables

```
rm(list=ls())
```

```
myFunc_04 <- function(x) x + 99
```

```
myFunc_05 <- function() {  
  myFunc_04 <- function(x) x * 2  
  myFunc_04(20)  
}
```

```
myFunc_05()
```

What does the preceding code return? **Key point:** **don't** give identical names to functions and variables

# New Functional Environments for Each Execution



- Every time a function is called, a new environment is called to host execution; each invocation is completely independent
- The following function returns a value of 999 every time

```
# NOTE  rm(list=ls()) is deleted
```

```
myFunc_06 <- function() {  
  if(!exists("myAtomicVector")){  
    myAtomicVector <- 999  
  } else {  
    myAtomicVector <- myAtomicVector + 1  
  }  
  print(myAtomicVector)  
}
```

```
myFunc_06()
```



A function will search for a value when it's run, **not** when it's created

```
> rm(list=ls())
```

```
> myFunc_07 <- function() x
```

```
> x <- 15
```

```
> myFunc_07()  
[1] 15
```

```
> x <- 20
```

```
> myFunc_07()  
[1] 20
```





- Variables internal to a function, i.e., variables which are not passed to a function, should be locally scoped to ensure that a function is self-contained
- A function that is not self-contained can cause a pernicious error that can be difficult to identify
- Use the `findGlobals` function from the `codetools` package to identify global variables in a function



```
> rm(list=ls())

> myFunc_08 <- function() x + 1

# NOTE myFunc_08 is not self-contained

> codetools::findGlobals(myFunc_08)
[1] "+" "x"
```



- Write any function with locally-scoped variables, confirming there are locally scoped using the `codetools` package



- It is important to distinguish between the formal and actual arguments of a function
- **Formal arguments** are a property of the function

## Arithmetic Mean

### Description

Generic function for the (trimmed) arithmetic mean.

### Usage

```
mean(x, ...)  
  
## Default S3 method:  
mean(x, trim = 0, na.rm = FALSE, ...)
```

### Arguments

- x** An R object. Currently there are methods for numeric/logical vectors and [date](#), [date-time](#) and [time interval](#) objects. Complex vectors are allowed for `trim = 0`, only.
- trim** the fraction (0 to 0.5) of observations to be trimmed from each end of `x` before the mean is computed. Values of `trim` outside that range are taken as the nearest endpoint.
- na.rm** a logical value indicating whether NA values should be stripped before the computation proceeds.
- ...** further arguments passed to or from other methods.



- It is important to distinguish between the formal and actual arguments of a function
- **Actual or calling arguments** can vary each time you call a function

```
> mean(x = 1:10)
[1] 5.5
```

```
> mean(x = 99:999)
[1] 549
```

- In the above examples, the calling arguments are `1:10` and `99:999` respectively



- When calling a function you can specify arguments by position, by complete name, or by partial name
- Arguments are matched in the following order
  - 1 Exact name (perfect matching)
  - 2 Prefix matching (imperfect/partial matching)
  - 3 Position

```
myFunc_09 <- function(arg1, my_arg2, my_arg3){  
  list(a = arg1, m1 = my_arg2, m2 = my_arg3)  
}
```

# Calling Arguments of a Function [CONT'D]



```
# positional                                # joint partial matching and positional
> str(myFunc_09(1, 2, 3))                    > str(myFunc_09(2, 3, a = 1))
List of 3                                    List of 3
 $ a : num 1                                $ a : num 1
 $ m1: num 2                                $ m1: num 2
 $ m2: num 3                                $ m2: num 3

# exact matching and positional
> str(myFunc_09(2, 3, arg1 = 1))
List of 3
 $ a : num 1
 $ m1: num 2
 $ m2: num 3
```



- You only want to use positional matching for the first one or two arguments of a function call, i.e., the most commonly used arguments
- Avoid using positional matching for infrequently used arguments
- If a function uses `...` (ellipsis), you can only specify arguments listed after the `...` with their full name, i.e., exact matching
- If you are writing code for a package to be published to CRAN, you are not permitted to use partial matching





- If you wish call a function with a list of arguments, use the following code

```
> funcArguments <- list(1:10, na.rm = TRUE)
```

```
> do.call(mean, funcArguments)
```

```
[1] 5.5
```

```
# equivalent to
```

```
> mean(1:10, na.rm = TRUE)
```

```
[1] 5.5
```



```
# w/o default values
myFunc_10 <- function(a, b) {
  c(a, b)
}
```

```
> myFunc_10()
```

```
Error in myFunc_10() : argument "a" is missing, with no default
```

```
# with default values
myFunc_11 <- function(a = 1, b = 2) {
  c(a, b)
}
```

```
> myFunc_11()
```

```
[1] 1 2
```



Function arguments in R can be defined in terms of other arguments

```
myFunc_12 <- function(a = 1, b = a * 2) {  
  c(a, b)  
}
```

```
> myFunc_12()  
[1] 1 2
```

```
> myFunc_12(111)  
[1] 111 222
```

```
> myFunc_12(99, 100)  
[1] 99 100
```



Two common approaches to determine whether or not an argument was supplied to a function:

## 1) `missing()`

```
myFunc_13 <- function(arg1, arg2) {  
  c(missing(arg1), missing(arg2))  
}
```

```
> myFunc_13()  
[1] TRUE TRUE
```

```
> myFunc_13(arg1 = 1)  
[1] FALSE TRUE
```

```
> myFunc_13(arg2 = 99)  
[1] TRUE FALSE
```



- 2) Set default argument values to `NULL` and subsequently test if the argument is supplied using `is.null()`

```
> myFunc_14 <- function(arg1 = NULL, arg2 = NULL) {  
  c(is.null(arg1), is.null(arg2))  
}
```

```
> myFunc_14()  
[1] TRUE TRUE
```

```
> myFunc_14(arg1 = 1)  
[1] FALSE TRUE
```

```
> myFunc_14(arg2 = 99)  
[1] TRUE FALSE
```

# Lazy Functional Evaluation of Calling Arguments



- R function arguments are only evaluated when they are used
- If you want to ensure that an argument is evaluated you can use `force()`

```
myFunc_15 <- function(x) {  
  10  
}
```

```
> myFunc_15()
```

```
[1] 10
```

```
> myFunc_15(thisIsNonsense)
```

```
[1] 10
```

```
> myFunc_15("nonsense")
```

```
[1] 10
```

```
myFunc_16 <- function(x) {  
  force(x)  
  10  
}
```

```
> myFunc_16(thisIsNonsense)
```

```
Error in force(x) : object  
  'thisIsNonsense' not found
```



- Default arguments are evaluated inside the function
- If the expression depends on the current environment the results will differ depending on whether you use the default value or explicitly provide one

```
myFunc_17 <- function(a = ls()){  
  z <- 10  
  a  
}
```

```
> myFunc_17()  
[1] "a" "z"
```

```
> myFunc_17(ls())  
[1] "i"          "j"          "myFunc_13"  
[4] "myFunc_15" "myFunc_17"
```



The last expression evaluated in a function becomes the return value

```
myFunc_18 <- function(xyz) {  
  if (xyz < 10) {  
    0  
  } else {  
    10  
  }  
}
```

```
> myFunc_18(5)  
[1] 0
```

```
> myFunc_18(10)  
[1] 10
```



# To `return()` or not to `return()`



- The last expression evaluated in a function is the return value
- You can always wrap the final expression in `return()` if you choose
- Calling `return()` is an additional call and will add to the execution time of your function, albeit minuscule for a single call
- In simplistic functions, R programmers will typically omit `return()`
- In longer, more complicated functions, `return()` is often used to distinguish “leaves” of code
- In sum, for the purposes of this class, I require the use of `return()` to make the code more legible for any functions with “leaves” of code

# To return() or not to return()



```
# simple function, does not require a return()
```

```
myFunc_15 <- function(x){  
  10  
}
```

```
# a more complex function benefits visually from having return()  
#   but does not require return()
```

```
myFunc_18 <- function(xyz) {  
  if (xyz < 10) {  
    return(0)  
  } else {  
    return(10)  
  }  
}
```



- 1 Write a function that takes two arguments, `firstRow` and `lastRow`, and returns rows `firstRow` through `lastRow` of `iris`, and subsequently call the function with values `firstRow = 1` and `lastRow = 3`, using both positional matching and exact matching
- 2 In the question above, what are the formal and calling arguments of the function?
- 3 Is this function self-contained? Why or why not?
- 4 Rewrite the above function to include a data frame `myDataFrame` as an additional argument, such that it returns rows `firstRow` through `lastRow` of `myDataFrame`
- 5 Rewrite the function to use default arguments `firstRow = 1` and `lastRow = 10`, and evaluate all 3 arguments at the beginning of the function using `force`



## Debugging

How to fix unanticipated problems

## Condition Handling

How functions communicate problems and how actions can be taken based on those communications

## Defensive Programming

How to avoid common problems before they occur



There are three key debugging tools


- 1 Error inspector and `traceback()` which lists a sequence of calls that lead to the error
- 2 “Return with Debug” tool and `options(error = browser)` which open an interactive session where the error occurred
- 3 Breakpoints and `browser()` which open an interactive session at an arbitrary location in the code



# A Brief Digression

Depending on your selection in the menu bar, different actions will occur when `R` throws an error

- Selecting **Message Only** will simply print an error message to the console
- Selecting **Error Inspector** additionally provides links to *Show Traceback* and *Rerun with Debug*
- Selecting **Break in Code** additionally launches *Browse on Error*





- The call stack is the sequence of calls that lead up to an error
- For example, if we run the following code...

```
> rm(list=ls())

> myFunc_21 <- function(arg_01) myFunc_22(arg_01)
> myFunc_22 <- function(arg_02) myFunc_23(arg_02)
> myFunc_23 <- function(arg_03) myFunc_24(arg_03)
> myFunc_24 <- function(arg_04) "myString" + arg_04
```

- ... and then call `myFunc_21()`, we see the following error message

```
> myFunc_21(99)
Error in "myString" + arg_04 : non-numeric argument to binary operator
```



- The call stack is the sequence of calls that lead up to an error
- For example, if we run the following code...

```
> rm(list=ls())

> myFunc_21 <- function(arg_01) myFunc_22(arg_01)
> myFunc_22 <- function(arg_02) myFunc_23(arg_02)
> myFunc_23 <- function(arg_03) myFunc_24(arg_03)
> myFunc_24 <- function(arg_04) "myString" + arg_04
```

- ... and then call `myFunc_21()`, we see the following error message

```
> myFunc_21(99)
Error in "myString" + arg_04 : non-numeric argument to binary operator
```





- Looking at the **Console** pane, you should see the following

```
> myFunc_21(99)
```

```
Error in "myString" + arg_04 : non-numeric argument to binary operator
```

Hide Traceback

Rerun with Debug

```
4 myFunc_24(arg_03)
3 myFunc_23(arg_02)
2 myFunc_22(arg_01)
1 myFunc_21(99)
```

- The call stack is to be read from bottom to top:
  - The initial call is to `myFunc_21()`
  - `myFunc_21()` calls `myFunc_22()`
  - `myFunc_22()` calls `myFunc_23()`
  - `myFunc_23()` calls `myFunc_24()` which triggers the error
- The **Traceback** window shows you where the error occurred, **not** why it occurred



- Selecting *Rerun with Debug* allows you to enter the interactive debugger
- This reruns the command that create the error, pausing the execution where the error occurred
- This puts you in an interactive state inside the function, and you can interact with any objects defined there
- You will observe
  - 1 A **Traceback** pane with the call stack
  - 2 An **Environment** pane with all objects in the current environment
  - 3 A **Code Browser** pane (icon of glasses) listing the statement that will be run next highlighted in yellow
  - 4 A `Browse[1] >` prompt in the console window which allows you to run arbitrary code

# Browsing on Error



The screenshot shows the RStudio interface with a menu bar at the top (RStudio, File, Edit, Code, View, Plots, Session, Build, Debug, Tools, Window, Help) and a status bar at the bottom (RStudio, 100%, Tue 12:19, Paul Intrevado). The console window on the left displays the following error message:

```
> myFunc_21(99)
Error in "myString" + arg_04 : non-numeric argument to binary operator
Called from: myFunc_24(arg_03)
Browse[1]>
```

The traceback window on the right shows the following stack trace:

```
1 function(arg_04) "myString" + arg_04
```

The Environment window shows the following values:

myFunc_240
arg_04
99


The Files pane at the bottom left shows the file explorer with the following structure:

- Home
  - .R
  - .RData
  - .Rhistory
  - Applications
  - Desktop
  - Documents
  - Downloads
  - Dropbox

# Browsing on Error



A few special commands can be accessed in the toolbar on the **Console** pane (from left to right)



- **Next** executes the next step in the function
- **Step Into** works similarly to **Next**, except if the next line is a function, it will also step into that function
- **Finish** completes execution of current loop or function
- **Continue** leaves interactive debugging and continues regular execution of the function
- **Stop** stops debugging, terminates function, and returns to the global workspace



- The task of handling expected errors, e.g., when your function is expecting an atomic vector as an argument but is passed a data frame
- In R, there are two main tools for handling conditions (including errors) programatically
  - 1 `try()` gives you the ability to continue execution even when an error occurs
  - 2 `tryCatch()` lets you specify *handler* functions that control what happens when a condition is signaled



Wrapping code in the statement `try()` results in an error message printing **but** execution will continue

```
> rm(list=ls())
```

```
myFunc_25 <- function(z){  
  log(z)  
  print("Made it here")  
}
```

```
> myFunc_25("abc")
```

```
Error in log(z) : non-numeric argument to mathematical function
```



```
myFunc_26 <- function(z){  
  try(log(z))  
  print("Made it here")  
}
```

```
> myFunc_26("abc")
```

```
Error in log(z) : non-numeric argument to mathematical function  
[1] "Made it here"
```

# Ignoring Errors with `try()`



- If you prefer, you can suppress the error message with `try(..., silent = TRUE)`
- The output of `try()` can also be captured
  - 1 If the execution of code within `try()` is successful, the result will be the last result evaluated (just as in a function)
  - 2 If unsuccessful, the (invisible) result will be of class `try-error`

```
> successful <- try(1 + 99)
```

```
> class(successful)
[1] "numeric"
```

```
> unsuccessful <- try("a" + "b")
```

```
Error in "a" + "b" : non-numeric argument to binary operator
```

```
> class(unsuccessful)
[1] "try-error"
```





- `tryCatch()` is a general tool for handling conditions
- `tryCatch()` can handle
  - 1 errors (made by `stop()`)
  - 2 warnings (`warning()`)
  - 3 message (`message()`)
  - 4 interrupts (user-terminated code execution, e.g., `ctrl + C`)
- `tryCatch()` maps conditions to **handlers**, i.e., named functions that are called with the condition as an argument
- If a condition is signaled, `tryCatch()` will call the first handler whose name matches one of the classes of the condition

# Handling Conditions with `tryCatch()`



```
show_condition <- function(code) {  
  tryCatch(code,  
            error = function(x) "myError",  
            warning = function(x) "myWarning",  
            message = function(x) "myMessage"  
  )  
}
```

```
> show_condition(stop("!"))  
[1] "myError"
```

```
> show_condition(warning("?"))  
[1] "myWarning"
```

```
> show_condition(message("?"))  
[1] "myMessage"
```

# Handling Conditions with `tryCatch()`



Let's follow the execution of the function `show_condition()`

- 1 `show_condition(stop("!"))` calls the function `show_condition()`, passing `stop("!")` as the argument, represented in the function as `code`
- 2 `code` is executed in the `tryCatch()` block, where `code == stop("!")`
- 3 the function `stop()` *“stops execution of the current expression and executes an error action”*
- 4 when `stop()` executes an error action, `tryCatch()` maps the **error** condition to a function `error = function(x)` `"myError"`, which prints the word `myError` to the console
- 5 execution of the function terminates

# Handling Conditions with `tryCatch()`



- When a condition is mapped to a function, what is being passed to that function?
- Let's modify the previous code and explore the inner workings of condition handling

```
show_condition <- function(code) {  
  tryCatch(code,  
            error = function(x) y <- x  
  )  
}
```

```
> show_condition(stop("!"))  
## this call generates no message in the console
```

- This is the first time we observe the `<-` operator, which makes an assignment to a *global* variable

# Handling Conditions with `tryCatch()`



```
> y
<simpleError in doTryCatch(return(expr), name, parentenv, handler): !>

> str(y)
List of 2
 $ message: chr "!"
 $ call    : language doTryCatch(return(expr), name, parentenv, handler)
 - attr(*, "class")= chr [1:3] "simpleError" "error" "condition"

> attributes(y)
$names
[1] "message" "call"

$class
[1] "simpleError" "error"      "condition"

> y$message
[1] "!"
```

# Handling Conditions with `tryCatch()`



- `tryCatch()` can be customized:

```
show_condition <- function(code) {  
  tryCatch(code,  
    error = function(x) {  
      print(x$message)  
      print(x$call)  
      writeLines("\nSilly error!")  
    }  
  )  
}  
  
> show_condition(stop("!"))  
[1] "!"  
doTryCatch(return(expr), name, parentenv, handler)  
  
Silly error!
```



Write a function employing error handling techniques that takes a single vector as input, take the natural log of each element in that vector, and print the result of each to the console



- Defensive programming is the art of making code fail in a well-defined manner even when something unexpected occurs
- A key principle of defensive programming is to *fail fast*: as soon as something wrong is discovered, signal an error
- This *fail fast* behavior is more work up front for the programmer, but results in easier debugging for the user, as they receive errors earlier rather than later, before the error has been potentially digested by multiple functions





- Be strict about what a function accepts
  - If a function is not vectorized in inputs but uses functions that are, build in a check to ensure that inputs are scalars
  - Use `stopifnot()` or the `assertthat` package
- Avoid functions that use non-standard evaluation such as `subset`, `transform` and `with`
  - These functions save time when working with R interactively, but they typically fail uninformatively
  - Non-standard evaluation is the ability of a computing language to access not only the value(s) of a function's argument but also the code used to compute them (Advanced R, Chapter 13)



- Be strict about what a function accepts
  - If a function is not vectorized in inputs but uses functions that are, build in a check to ensure that inputs are scalars
  - Use `stopifnot()` or the `assertthat` package
- Avoid functions that use non-standard evaluation such as `subset`, `transform` and `with`
  - These functions save time when working with R interactively, but they typically fail uninformatively
  - Non-standard evaluation is the ability of a computing language to access not only the value(s) of a function's argument but also the code used to compute them (Advanced R, Chapter 13)



- Avoid functions that return different of output depending on their input
  - Two big offenders are `[` and `sapply()`
    - Whenever subsetting a data frame in a function, **always** use the option `drop = F` to maintain the data structure, e.g., to avoid converting a one-column data frame to an atomic vector



# stop() versus stopifnot()

```
> myVec <- c("a", "bcd", "efgh")
```

```
> if(length(unique(nchar(myVec))) != 1) {  
  stop("Error: Elements of your input vector do not have the  
  same length!")  
}
```

Error: Error: Elements of your input vector do not have the same length!

```
> stopifnot(length(unique(nchar(myVec))) != 1,  
  "Error: Elements of your input vector HAVE the same length!")  
Error: "Error: Elements of your input vector HAVE the same length!"  
is not TRUE
```

```
> stopifnot(1 == 1, all.equal(pi, 3.14159265), 1 < 2) # all TRUE
```

```
> stopifnot(1 == 2, all.equal(pi, 3.14159265), 1 < 2) # all first  
#is FALSE
```

```
Error: 1 == 2 is not TRUE
```